MATH3431 Machine Learning and Neural Networks III — Epiphany term 2025

Exercise sheet

Lecturer/Author: Georgios P. Karagiannis — georgios.karagiannis@durham.ac.uk

Part 1. Convex learning problems

Exercise 1. ($\star$)Let $f:\mathbb{R}^d\to\mathbb{R}$ such that $f(w)=g(<w,x>+y)$ or some $x\in\mathbb{R}^d$, $y\in\mathbb{R}$. Show that: If g is convex function then $f$is convex function.

Solution. Let $u,v\in\mathbb{R}^d$ and $a\in[0,1]$. It is

$$\begin{aligned}
f(\alpha u+(1-\alpha)v) =& g(<\alpha u+(1-\alpha)v,x>+y)\\
=& g(<\alpha u,x>+<(1-\alpha)v,x>+y)\\
=& g(\alpha(<u,x>+y)+(1-\alpha)(<v,x>+y)) \qquad y=\alpha y+(1-\alpha)y\\
\leq& \alpha g(<u,x>+y)+(1-\alpha)g(<v,x>+y) \qquad (g \text{ is convex})\\
=& \alpha f(u)+(1-\alpha)f(v)
\end{aligned}$$

Exercise 2. ($\star$)Let functions g_1 be ρ_1-Lipschitz and g_2 be ρ_2-Lipschitz. Then, show that, f with $f(x)=g_1(g_2(x))$ is $\rho_1\rho_2$-Lipschitz.

Solution.

$$\begin{aligned}
|f(w_1)-f(w_2)| =& |g_1(g_2(w_1))-g_1(g_2(w_2))|\\
\leq& \rho_1|g_2(w_1)-g_2(w_2)|\\
\leq& \rho_1\rho_2|w_1-w_2|
\end{aligned}$$

Exercise 3. ($\star$)Let $f:\mathbb{R}^d\to\mathbb{R}$ with $f(w)=g(\langle w,x\rangle+y)$ $x\in\mathbb{R}^d$ and $y\in\mathbb{R}$. Let $g:\mathbb{R}\to\mathbb{R}$ be a β-smooth function. Then show that f is a $\left(\beta\|x\|^2\right)$-smooth.

Hint:: You may use Cauchy-Schwarz inequality $\langle y,x\rangle\leq\|y\|\,\|x\|$

$$
\begin{aligned}
f(v) =& g(\langle w, x\rangle + y) \\
\leq& g(\langle w, x\rangle + y) + g'(\langle w, x\rangle + y)\langle v - w, x\rangle + \frac{\beta}{2}(\langle v - w, x\rangle)^2 && (g \text{ is smooth}) \\
\leq& g(\langle w, x\rangle + y) + g'(\langle w, x\rangle + y)\langle v - w, x\rangle + \frac{\beta}{2}(\|v - w\|\|x\|)^2 && (\text{Cauchy-Schwatz inequality}) \\
=& f(w) + \langle \nabla f(w), v - w\rangle + \frac{\beta\|x\|^2}{2}\|v - w\|^2
\end{aligned}
$$

Exercise 4. ($\star$)Show that $f : S \to \mathbb{R}$ is ρ-Lipschitz over an open convex set S if and only if for all $w \in S$ and $v \in \partial f(w)$ it is $\|v\| \leq \rho$.

Hint:: You may use Cauchy-Schwarz inequality $\langle y, x\rangle \leq \|y\|\|x\|$

Solution. $\Longrightarrow$ Let $f : S \to \mathbb{R}$ be ρ-Lipschitz over convex set S, $w \in S$ and $v \in \partial f(w)$.

- Since S is open we get that there exist $\epsilon > 0$ such as $u := w + \epsilon\frac{v}{\|v\|}$ where $u \in S$. So $\langle u - w, v\rangle = \epsilon\|v\|$ and $\|u - w\| = \epsilon$.
- From the subgradient definition we get
$$f(u) - f(w) \geq \langle u - w, v\rangle = \epsilon\|v\|$$
- From the Lipschitzness of $f(\cdot)$ we get
$$f(u) - f(w) \leq \rho\|u - w\| = \rho\epsilon$$
Therefore $\|v\| \leq \rho$.
$\Longleftarrow$ It is for all $w \in S$ and $v \in \partial f(w)$ it is $\|v\| \leq \rho$.
- For any $u \in S$, it is
$$
\begin{aligned}
f(w) - f(u) \leq& \langle v, w - u\rangle && (\text{ because } v \in \partial f(w)\ ) \\
\leq& \|v\|\|w - u\| && \text{by Cauchy-Schwarz inequality} \\
\leq& \rho\|w - u\| && \text{because } \|v\| \leq \rho
\end{aligned} \tag{0.1}
$$
- Similarly it results $u, w \in S$
$$f(w) - f(u) \leq \langle v, u - w\rangle\|v\| \leq \|v\|\|u - w\| \leq \rho\|u - w\|$$
from (0.1) because w, u can be swaped in (0.1) as they both are any values in S.

Exercise 5. ($\star$)Let $g_1(w), ..., g_r(w)$ be r convex functions, and let $f(\cdot) = \max_{\forall j}(g_j(\cdot))$. Show that for some w it is $\nabla g_k(w) \in \partial f(w)$ where $k = \arg\max_j(g_j(w))$ is the index of function $g_j(\cdot)$ presenting the greatest value at w.

Solution. Since g_k is convex, for all u

$$g_k(u) \geq g_k(w) + \langle u - w, \nabla g_k(w) \rangle$$

However $f(u) = \max_{\forall j}(g_j(u)) \geq g_k(u)$ for any j, and $f(w) = g_k(w)$ at w. Then

$$\begin{aligned} f(u) \geq & g_k(u) \\ \geq & g_k(w) + \langle u - w, \nabla g_k(w) \rangle \\ = & f(w) + \langle u - w, \nabla g_k(w) \rangle \end{aligned}$$

Then by the definition of the sub-gradient $\nabla g_k(w) \in \partial f(w)$

Exercise 6. ($\star$)Consider the regression learning problem $(\mathcal{H}, \mathcal{Z}, \ell)$ with predictor rule $h(x) = \langle w, x \rangle$ labeled by some unknown parameter $w \in \mathcal{W}$, loss function $\ell(w, (x, y)) = (\langle w, x \rangle - y)^2$, feature $x \in \mathcal{X}$, and target $y \in \mathbb{R}$. Let $\mathcal{W} = \mathcal{X} = \left\{\omega \in \mathbb{R}^d : |\omega| \leq \rho\right\}$ for some $\rho > 0$.

(1) Show that the resulting learning problem is Convex-Lipschitz-Bounded learning problem.
(2) Specify the parameters of Lipschitnzess.

Solution. According to the definitions given in the lecture:

- Convex-Lipschitz-Bounded Learning Problem $(\mathcal{H}, \mathcal{Z}, \ell)$ with parameters ρ, and B, is called the learning problem whose the hypothesis class $\mathcal{H}$ is a convex set, for all $w \in \mathcal{H}$ it is $\|w\| \leq B$, and the loss function $\ell(\cdot, z)$ is convex and ρ-Lipschitz function for all $z \in \mathcal{Z}$.

I have:

Convexity: The function $g : \mathbb{R} \to \mathbb{R}$, defined by $g(a) = a^2$ is convex convex. Eg. $\frac{\mathrm{d}^2}{\mathrm{d}a^2} g(a) = 1 \geq 0$ is non-negative. The convexity of $\ell(w, z = (x, y))$ for all z follows as a composition of g with a linear function.

Lipschitzness: The function $g(a) = a^2$ is 1-Lipschitz since It is

$$|g(a_2) - g(a_1)| = \left|a_2^2 - a_1^2\right| = |(a_2 + a_1)(a_2 - a_1)| \leq 2\rho(a_2 - a_1) = 2\rho|a_2 - a_1|$$

Hence because $|x| \leq \rho$, $g(a)$ is $2\rho^2$-Lipschitz as a composition.

Boundness: The norm of each hypothesis w is bounded by ρ according to the assumptions.

Therefore,

(1) the learning problem under consideration is a Convex-Lipschitz-Bounded learning problem.
(2) the parameter of Lipschitzness is $2\rho^2$.

Exercise 7. ($\star$)If $f$is λ-strongly convex and u is a minimizer of f then for any w

$$f(w) - f(u) \geq \frac{\lambda}{2} \|w - u\|^2$$

Hint:: Use the definition, and set $\alpha \to 0$.

Solution.

Exercise 8. ($\star$) Let $f : \mathbb{R}^d \to \mathbb{R}$ be a convex and β-smooth function.

(1) Show that for $v, w \in \mathbb{R}^d$

$$f(v) - f(w) \in \left(\langle \nabla f(w), v - w \rangle, \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|^2 \right)$$

(2) Show that for $v, w \in \mathbb{R}^d$ such that $v = w - \frac{1}{\beta} \nabla f(w)$, it is

$$\frac{1}{2\beta} \|\nabla f(w)\|^2 \leq f(w) - f(v)$$

(3) Additionally assume that $f(x) > 0$ for all $x \in \mathbb{R}^d$. Show that for $w \in \mathbb{R}^d$,

$$\|\nabla f(w)\| \leq \sqrt{2\beta f(w)}$$

Solution.

Exercise 9. ($\star$)Let $f : \mathbb{R}^d \to \mathbb{R}$ be a λ-strongly convex function. Assume that w^* is a minimizer of f i.e.

$$w^* = \arg\min_w \{f(w)\}$$

Show that for any $w \in \mathbb{R}^d$ it holds

$$f(w) - f(w^*) \geq \frac{\lambda}{2} \|w - w^*\|^2$$

Hint: Use the definition of λ-strongly convex function, properly rearrange it, and ...

Solution.

Exercise 10. ($\star$)Show that the function $J(x; \lambda) = \lambda \|x\|^2$ is 2λ-strongly convex

Solution. We just need to check that for all w, u, and $\alpha \in (0, 1)$ we have

$$J(aw + (1-\alpha)u; \lambda) \leq aJ(w; \lambda) + (1-\alpha) J(u; \lambda) - \frac{2\lambda}{2} \alpha(1-\alpha) \|w - u\|^2 \iff$$
$$\|aw + (1-\alpha)u\|_2^2 \leq a\|w\|_2^2 + (1-\alpha)\|u\|_2^2 - a(1-\alpha)\|w - u\|_2^2 \iff 0 \leq 0$$

Exercise 11. ($\star\star\star$)Consider a learning problem $(\mathcal{H}, \mathcal{Z}, \ell)$ with $\mathcal{H} \subset \mathbb{R}^d$, $d > 0$, and loss function $\ell : \mathcal{H} \times \mathcal{Z} \to \mathbb{R}_+$ which is convex, β-smooth and non-negative. Let $\mathfrak{A}$ be a learning algorithm with output $\mathfrak{A}(\mathcal{S})$ trained against training dataset $\mathcal{S} = \{z_1, ..., z_m\}$ of IID samples $z_1, ..., z_m \sim g$

where g is a data generating distribution. In particular, consider that $\mathfrak{A}(\mathcal{S})$ is the Regularized Loss Minimization learning rule that outputs a hypothesis in

$$\min_{w} \left\{ \hat{R}_{\mathcal{S}}(w) + \lambda \|w\|_2^2 \right\}$$

for $\lambda \geq \frac{2\beta}{m}$ where $\hat{R}_{\mathcal{S}}(w) = \frac{1}{m}\sum_{i=1}^{m} \ell(w, z_i)$ for all $w \in \mathcal{H}$.

(1) Prove that

$$\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) \right) \leq R_g(w) + \lambda \|w\|_2^2$$

for all $w \in \mathcal{H}$. $R_g(\cdot)$ denotes the risk function under the real data generating distribution g.

(2) Prove that

$$\mathrm{E}_{\mathcal{S}\sim g}\left(R_g(\mathfrak{A}(\mathcal{S})) - \hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) \right) \leq \frac{48\beta}{\lambda m} \mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) \right).$$

Hint:: If needed you can use the following:

Let $\mathcal{S}^{(i)} = \{z_1, ..., z_{i-1}, z', z_{i+1}, ..., z_m\}$ be a set resulting from $\mathcal{S}$ by replacing its i-th element z_i with an independently drawn $z' \sim g$. Then

$$24\beta\ell(\mathfrak{A}(\mathcal{S}), z_i) + \lambda m \ell(\mathfrak{A}(\mathcal{S}), z_i) + 24\beta\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right), z'\right) - \lambda m \ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right), z_i\right) \geq 0$$

(3) Show that the learning algorithm $\mathfrak{A}$ is on-average-replace-one-stable with rate ε. Specify that rate ε as a function of β, λ, m and possibly any other user specified constants if needed. Explain how the shrinkage parameter λ, the training dataset size m, and the smoothness parameter β affect the stability of the learning algorithm $\mathfrak{A}$.

(4) Show that the expected risk is bounded as follows

$$\mathrm{E}_{\mathcal{S}\sim g}\left(R_g(\mathfrak{A}(\mathcal{S})) \right) \leq \left(1 + \frac{48\beta}{\lambda m}\right)\left(R_g(w) + \lambda \|w\|_2^2 \right)$$

for all $w \in \mathcal{H}$.

Solution.

(1) We have

$$\begin{aligned} \hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) &\leq \hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) + \lambda \|\mathfrak{A}(\mathcal{S})\|_2^2 \\ &\leq \hat{R}_{\mathcal{S}}(w) + \lambda \|w\|_2^2; \quad \forall w \in \mathcal{W} \end{aligned}$$

and by taking expectations w.r.t. $\mathcal{S}$, it is

$$\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) \right) \leq R_g(w) + \lambda \|w\|_2^2; \quad \forall w \in \mathcal{W} \tag{0.2}$$

because $\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}(\cdot) \right) = \frac{1}{m}\sum_{i=1}^{m} \mathrm{E}_{\mathcal{S}\sim g}(\ell(\cdot, z_i)) = R_g(\cdot)$.

(2) From a well known theorem to us, it is

$$\mathrm{E}_{\mathcal{S}\sim g}\left(R_g(\mathfrak{A}(\mathcal{S})) - \hat{R}_{\mathcal{S}}(\mathfrak{A}(\mathcal{S})) \right) = \mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right), z_i\right) - \ell(\mathfrak{A}(\mathcal{S}), z_i) \right)$$

Now I ll gonna work on the second term as this is what I ve given in the hint...

It is

$$24\beta\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)+\lambda m\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)+24\beta\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z'\right)-\lambda m\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)\geq 0\Leftrightarrow$$

$$\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)-\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\leq\frac{24\beta}{\lambda m}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)+\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z'\right)\right)$$

Taking expectations

$$\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)-\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\right)\leq\frac{24\beta}{\lambda m}\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)+\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z'\right)\right)$$

Due to the sampling it is

$$\mathrm{E}_{\mathcal{S}}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\right)=\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z'\right)\right)=\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)$$

So I get

$$\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)-\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\right)\leq\frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right).$$

So I get

$$\mathrm{E}_{\mathcal{S}\sim g}\left(R_g\left(\mathfrak{A}\left(\mathcal{S}\right)\right)-\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)\leq\frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right).$$

(3) Okay, let's say, I did not do the previous part. I see it is

$$\mathrm{E}_{\mathcal{S}\sim g}\left(R_g\left(\mathfrak{A}\left(\mathcal{S}\right)\right)-\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)\leq\frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right).$$

From a well known theorem to us, it is

$$\mathrm{E}_{\mathcal{S}\sim g}\left(R_g\left(\mathfrak{A}\left(\mathcal{S}\right)\right)-\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)=\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)-\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\right)$$

So

$$\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)-\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\right)\leq\frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)$$

but the expectation depends on m... so

$$\begin{aligned}
\mathrm{E}_{\mathcal{S},z',i}\left(\ell\left(\mathfrak{A}\left(\mathcal{S}^{(i)}\right),z_i\right)-\ell\left(\mathfrak{A}\left(\mathcal{S}\right),z_i\right)\right) &\leq \frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)\\
&\leq \frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(0\right)\right)\\
&\leq \frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\max\hat{R}_{\mathcal{S}}\left(0\right)\right)\\
&\leq \frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\frac{1}{m}\sum_{i=1}^{n}\underbrace{\max\ell\left(0,z_i\right)}_{=C}\right)\\
&\leq \frac{48\beta}{\lambda m}C
\end{aligned}$$

so it is the learning algorithm $\mathfrak{A}$ is on-average-replace-one-stable with rate

$$\varepsilon=\frac{48\beta}{\lambda m}C$$

and $C=\max\ell\left(0,z_i\right)$ ...or whatever constant they pick.

Larger training sample size m, and larger regularization parameter λ (eg more parsimonious model) lead to a more stable learning algorithm. Smaller smoothness parameter (the gradient changes less wrt the argument) leads to more stable learning algorithm.

(4) We use the decomposition discussed in the lectures,

$$\begin{aligned}
\mathrm{E}_{\mathcal{S}\sim g}\left(R_g\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right) &= \underbrace{\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)}+\underbrace{\mathrm{E}_{\mathcal{S}\sim g}\left(R_g\left(\mathfrak{A}\left(\mathcal{S}\right)\right)-\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)} \\
&\leq \mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right)+\frac{48\beta}{\lambda m}\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right) \\
&= \left(1+\frac{48\beta}{\lambda m}\right)\mathrm{E}_{\mathcal{S}\sim g}\left(\hat{R}_{\mathcal{S}}\left(\mathfrak{A}\left(\mathcal{S}\right)\right)\right) \\
&\leq \left(1+\frac{48\beta}{\lambda m}\right)\left(R_g\left(w\right)+\lambda\left\|w\right\|_2^2\right),\ \forall w\in\mathcal{H}
\end{aligned}$$

Part 2. Stochastic learning

Exercise 12. ($\star$) Assume a Bayesian model

$$\begin{cases} z_i|w & \overset{\text{ind}}{\sim} f\left(z_i|w\right),\ i=1,...,n \\ w & \sim f\left(w\right) \end{cases}$$

and consider that our objective is the discovery of MAP estimate w^* i.e.

$$w^*=\arg\min_{\forall w\in\Theta}\left(-\log\left(L_n\left(w\right)\right)-f\left(w\right)\right)=\arg\min_{\forall w\in\Theta}\left(-\sum_{i=1}^{n}\log\left(f\left(z_i|w\right)\right)-\log\left(f\left(w\right)\right)\right)$$

by using SGD with update

$$w^{(t+1)}=w^{(t)}+\eta_t\left(\frac{n}{m}\sum_{j\in\mathcal{J}^{(t)}}\nabla_w\log\left(f\left(z_j|w^{(t)}\right)\right)+\nabla_w\log\left(f\left(w^{(t)}\right)\right)\right)$$

for some randomly selected set $\mathcal{J}^{(t)}\subseteq\{1,...,n\}^m$ of m integers from 1 to n via simple random sampling (SRS) with replacement. Show that

$$\mathrm{E}_{\mathcal{J}^{(t)}\sim\text{simple-random-sampling}}\left(\frac{n}{m}\sum_{j\in\mathcal{J}^{(t)}}\nabla_w\log\left(f\left(z_j|w^{(t)}\right)\right)\right)=\sum_{i=1}^{n}\nabla_w\log\left(f\left(z_i|w^{(t)}\right)\right)$$

Solution. It is

$$\begin{aligned}
\mathrm{E}_{\mathcal{J}^{(t)}\sim \mathrm{SRS}}\left(\frac{n}{m}\sum_{j\in\mathcal{J}^{(t)}}\nabla_w \log\left(f\left(z_j|w^{(t)}\right)\right)\right) =& \frac{n}{m}\sum_{j\in\mathcal{J}^{(t)}}\mathrm{E}_{\mathcal{J}^{(t)}\sim \mathrm{SRS}}\left(\nabla_w \log\left(f\left(z_j|w^{(t)}\right)\right)\right)\\
=& \frac{n}{m}\sum_{j\in\mathcal{J}^{(t)}}\mathrm{E}_{\mathcal{J}^{(t)}\sim \mathrm{SRS}}\left(\nabla_w \log\left(f\left(z_j|w^{(t)}\right)\right)\right)\\
=& \frac{n}{m}\sum_{j\in\mathcal{J}^{(t)}}\frac{1}{n}\sum_{i=1}^{n}\nabla_w \log\left(f\left(z_i|w^{(t)}\right)\right)\\
=& \sum_{i=1}^{n}\nabla_w \log\left(f\left(z_i|w^{(t)}\right)\right)
\end{aligned}$$

It is $\mathrm{E}_{\mathcal{J}^{(t)}\sim \mathrm{SRS}}\left(\nabla_w \log\left(f\left(z_j|w^{(t)}\right)\right)\right) = \frac{1}{n}\sum_{i=1}^{n}\nabla_w \log\left(f\left(z_i|w^{(t)}\right)\right)$ because the expectation is under the probability I get randomly an integer and for the jth on the probability is $1/n$ due to the random scheme. Also $\left|\mathcal{J}^{(t)}\right| = m$.

Exercise 13. ($\star$) Let $\{v_t; t = 1, ..., T\}$ be a sequence of vectors with $v_t \in \mathbb{R}^d$ and $d \in \mathbb{N} - \{0\}$. Consider an algorithm producing $\left\{w^{(t)}; t = 1, 2, 3, ...\right\}$ with

$$\begin{aligned}
w^{(1)} &= 0\\
w^{(t+1)} &= w^{(t)} - \eta v_t
\end{aligned}$$

$w_t \in \mathbb{R}^d$ and $d \in \mathbb{N} - \{0\}$. Show that

(1) it is

$$\langle w^{(t)} - w^*, v_t\rangle = \frac{1}{2\eta}\left(-\left\|w^{(t+1)} - w^*\right\|^2 + \left\|w^{(t)} - w^*\right\|^2\right) + \frac{\eta}{2}\left\|v_t\right\|^2$$

Hint:: Recall that

$$\|x + y\|_2^2 = \|x\|_2^2 + \|y\|_2^2 + 2\langle x, y\rangle,\ \forall x, y \in \mathbb{R}^d, d \in \mathbb{N} - \{0\}$$

(2) it is

$$\sum_{t=1}^{T}\langle w^{(t)} - w^*, v_t\rangle = \frac{1}{2\eta}\sum_{t=1}^{T}\left(-\left\|w^{(t+1)} - w^*\right\|^2 + \left\|w^{(t)} - w^*\right\|^2\right) + \frac{\eta}{2}\sum_{t=1}^{T}\left\|v_t\right\|^2$$

(3) (continue) it is

$$\sum_{t=1}^{T}\langle w^{(t)} - w^*, v_t\rangle \leq \frac{\left\|w^*\right\|^2}{2\eta} + \frac{\eta}{2}\sum_{t=1}^{T}\left\|v_t\right\|^2$$

Solution.

(1) It is

$$\begin{aligned}\left\langle w^{(t)}-w^*, v_t\right\rangle =& \frac{1}{\eta}\left\langle w^{(t)}-w^*, \eta v_t\right\rangle \\ =& \frac{1}{\eta}\left(-\left\langle w^{(t)}-w^*, -\eta v_t\right\rangle\right)\end{aligned}$$

Then by using the Hint as

$$\langle x, y\rangle = \frac{1}{2}\left(\|x+y\|_2^2 - \|x\|_2^2 - \|y\|_2^2\right)$$

for $x = w^{(t)} - w^* \in \mathbb{R}^d$ and $y = -\eta v_t \in \mathbb{R}^d$, I get

$$\begin{aligned}\left\langle w^{(t)}-w^*, v_t\right\rangle =& \frac{1}{2\eta}\left(-\left\|w^{(t)}-w^*-\eta v_t\right\|^2 + \left\|w^{(t)}-w^*\right\|^2 + \|-\eta v_t\|^2\right) \\ =& \frac{1}{2\eta}\left(-\left\|w^{(t+1)}-w^*\right\|^2 + \left\|w^{(t)}-w^*\right\|^2 + \eta^2\|v_t\|^2\right) \\ =& \frac{1}{2\eta}\left(-\left\|w^{(t+1)}-w^*\right\|^2 + \left\|w^{(t)}-w^*\right\|^2\right) + \frac{\eta}{2}\|v_t\|^2\end{aligned}$$

(2) So

$$\begin{aligned}\sum_{t=1}^{T}\left\langle w^{(t)}-w^*, v_t\right\rangle &= \frac{1}{2\eta}\sum_{t=1}^{T}\left(-\left\|w^{(t+1)}-w^*\right\|^2 + \left\|w^{(t)}-w^*\right\|^2\right) + \frac{\eta}{2}\sum_{t=1}^{T}\|v_t\|^2 \\ &= \frac{1}{2\eta}\left(\left\|w^{(1)}-w^*\right\|^2 - \left\|w^{(T+1)}-w^*\right\|^2\right) + \frac{\eta}{2}\sum_{t=1}^{T}\|v_t\|^2\end{aligned}$$

(3) So

$$\begin{aligned}\sum_{t=1}^{T}\left\langle w^{(t)}-w^*, v_t\right\rangle &= \frac{1}{2\eta}\left(\left\|w^{(1)}-w^*\right\|^2 - \left\|w^{(T+1)}-w^*\right\|^2\right) + \frac{\eta}{2}\sum_{t=1}^{T}\|v_t\|^2 \\ &\leq \frac{1}{2\eta}\left\|w^{(1)}-w^*\right\|^2 + \frac{\eta}{2}\sum_{t=1}^{T}\|v_t\|^2 \\ &= \frac{1}{2\eta}\|w^*\|^2 + \frac{\eta}{2}\sum_{t=1}^{T}\|v_t\|^2\end{aligned}$$

Exercise 14. ($\star$) Let $\{v_t; t = 1, ..., T\}$ be a sequence of vectors. Consider an algorithm producing $\left\{w^{(t)}; t = 1, 2, 3, ...\right\}$ with

$$\begin{aligned}w^{(1)} &= 0 \\ w^{\left(t+\frac{1}{2}\right)} &= w^{(t)} - \eta v_t \\ w^{(t+1)} &= \arg\min_{w\in\mathcal{H}}\left(\left\|w - w^{\left(t+\frac{1}{2}\right)}\right\|\right)\end{aligned}$$

for $t = 1, ..., T$.

Hint: You can use the following Lemma

(Projection Lemma): Let $\mathcal{H}$ be a closed convex set and let v be the projection of w onto $\mathcal{H}$,i.e.

$$v = \arg\min_{x \in \mathcal{H}} \|x - w\|^2$$

then for every $u \in \mathcal{H}$ it is

$$\|v - u\|^2 \leq \|w - u\|^2$$

Show that

(1) it is

$$\langle w^{(t)} - w^*, v_t \rangle \leq \frac{1}{2\eta} \left(- \left\| w^{(t+1)} - w^* \right\|^2 + \left\| w^{(t)} - w^* \right\|^2 \right) + \frac{\eta}{2} \|v_t\|^2$$

(2) it is

$$\sum_{t=1}^{T} \langle w^{(t)} - w^*, v_t \rangle \leq \frac{1}{2\eta} \sum_{t=1}^{T} \left(- \left\| w^{(t+1)} - w^* \right\|^2 + \left\| w^{(t)} - w^* \right\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^{T} \|v_t\|^2$$

(3) (continue) it is

$$\sum_{t=1}^{T} \langle w^{(t)} - w^*, v_t \rangle \leq \frac{\|w^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^{T} \|v_t\|^2$$

Comment: Above we show that Lemma **??** from "Handout **??**: Gradient descent" holds even when a projection step is included. Hence, even if a projection step is included after the update step of the recursion of GD algorithm or the SGD algorithm the analysis in Section **??** in "Handout **??**: Gradient descent" holds. Hence, even if a projection step is included after the update step of the recursion of SGD algorithm or the SGD algorithm the analysis in Section **??** in "Handout **??**: Stochastic gradient descent" holds.

Solution.

(1) It is

$$\begin{aligned}
\left\langle w^{(t)} - w^*, v_t \right\rangle &= \frac{1}{\eta} \left\langle w^{(t)} - w^*, \eta v_t \right\rangle \\
&= \frac{1}{2\eta} \left(- \left\| w^{(t)} - w^* - \eta v_t \right\|^2 + \left\| w^{(t)} - w^* \right\| + \eta^2 \|v_t\|^2 \right) \\
&= \frac{1}{2\eta} \left(- \left\| w^{(t+\frac{1}{2})} - w^* \right\|^2 + \left\| w^{(t)} - w^* \right\| + \eta^2 \|v_t\|^2 \right) \\
&= \frac{1}{2\eta} \left(- \left\| w^{(t+\frac{1}{2})} - w^* \right\|^2 + \left\| w^{(t)} - w^* \right\| \right) + \frac{\eta}{2} \|v_t\|^2 \\
&\leq \frac{1}{2\eta} \left(- \left\| w^{(t+1)} - w^* \right\|^2 + \left\| w^{(t)} - w^* \right\| \right) + \frac{\eta}{2} \|v_t\|^2
\end{aligned}$$

because from the Projection Lemma

$$\left\| w^{(t+1)} - w^* \right\|^2 \leq \left\| w^{(t+\frac{1}{2})} - w^* \right\|^2$$

(2) So

$$\sum_{t=1}^{T}\left\langle w^{(t)}-w^{*},v_{t}\right\rangle \leq\frac{1}{2\eta}\sum_{t=1}^{T}\left(-\left\|w^{(t+1)}-w^{*}\right\|^{2}+\left\|w^{(t)}-w^{*}\right\|\right)+\frac{\eta}{2}\sum_{t=1}^{T}\left\|v_{t}\right\|^{2}$$
$$=\frac{1}{2\eta}\left(\left\|w^{(1)}-w^{*}\right\|^{2}-\left\|w^{(T+1)}-w^{*}\right\|^{2}\right)+\frac{\eta}{2}\sum_{t=1}^{T}\left\|v_{t}\right\|^{2}$$

(3) So

$$\sum_{t=1}^{T}\left\langle w^{(t)}-w^{*},v_{t}\right\rangle \leq\frac{1}{2\eta}\left(\left\|w^{(1)}-w^{*}\right\|^{2}-\left\|w^{(T+1)}-w^{*}\right\|^{2}\right)+\frac{\eta}{2}\sum_{t=1}^{T}\left\|v_{t}\right\|^{2}$$
$$\leq\frac{1}{2\eta}\left\|w^{(1)}-w^{*}\right\|^{2}+\frac{\eta}{2}\sum_{t=1}^{T}\left\|v_{t}\right\|^{2}$$
$$=\frac{1}{2\eta}\left\|w^{*}\right\|^{2}+\frac{\eta}{2}\sum_{t=1}^{T}\left\|v_{t}\right\|^{2}$$

Exercise 15. ($\star$) [1]Consider the binary classification problem with inputs $x\in\mathcal{X}$ where $\mathcal{X}:=\left\{x\in\mathbb{R}^{d}:\left\|x\right\|_{2}\leq L\right\}$ for some given value $L>0$, target $y\in\mathcal{Y}$ where $\mathcal{Y}:=\{-1,+1\}$, and prediction rule $h_{w}:\mathbb{R}^{d}\rightarrow\{-1,+1\}$ with

$$h_{w}\left(x\right)=\text{sign}\left(w^{\top}x\right) \tag{0.3}$$
$$=\text{sign}\left(\sum_{j=1}^{d}w_{j}x_{j}\right) \tag{0.4}$$

Let the hypothesis class is

$$\mathcal{H}=\left\{x\rightarrow w^{\top}x:\forall w\in\mathbb{R}^{d}\right\} \tag{0.5}$$

In other words, the hypothesis $h_{w}\in\mathcal{H}$ is parametrized by $w\in\mathbb{R}^{d}$, it receives an input vector $x\in\mathcal{X}:=\mathbb{R}^{d}$ and it returns the label $y=\text{sign}\left(w^{\top}x\right)\in\mathcal{Y}:=\{\pm1\}$ where

$$\text{sign}\left(\xi\right)=\begin{cases}-1, & \text{if}\,\xi<0\\ +1, & \text{if}\,\xi>0\end{cases}$$

Consider a loss function $\ell:\mathbb{R}^{d}\rightarrow\mathbb{R}_{+}$ with

$$\ell\left(w,z=(x,y)\right)=\max\left(0,1-yw^{\top}x\right)+\lambda\left\|w\right\|_{2}^{2} \tag{0.6}$$

[1]We use standard notation

$$\text{sign}\left(\xi\right)=\begin{cases}-1, & \text{if}\,\xi<0\\ +1, & \text{if}\,\xi>0\end{cases}$$

±1 means either -1 or $+1$, $\mathbb{R}_{+}:=(0,+\infty)$, and $\left\|x\right\|_{2}:=\sqrt{\sum_{\forall j}\left(x_{j}\right)^{2}}$ for the Euclidean distance.

for some given value $\lambda > 0$.

Assume there is available a dataset of examples $S_n = \{z_i = (x_i, y_i)\,;\, i = 1, ..., n\}$ of size n.

Do the following:

(1) Show that the function $f : \mathbb{R} \to \mathbb{R}_+$ with $f(x) = \max(0, 1 - x)$ is convex in $\mathbb{R}$; and show that the loss (0.6) is convex.

Hint:: You may use Proposition **??** from Handout **??**: Elements of convex learning problems.

(2) Show that the loss $\ell(w, z)$ for $\lambda = 0$ (0.6) is L-Lipschitz (with respect to w) when $x \in \mathcal{X}$ where $\mathcal{X} := \left\{x \in \mathbb{R}^d : \|x\|_2 \leq L\right\}$.

Hint:: You may use the definition of Lipschitz function. Without loss of generality, you can consider any $w_1 \in \mathbb{R}^d$ and $w_2 \in \mathbb{R}^d$ such that $1 - yw_2^\top x \leq 1 - yw_1^\top x$, and then take cases $1 - yw_2^\top x > \text{or} < 0$ and $1 - yw_1^\top x > \text{or} < 0$ to deal with the max.

(3) Construct the set of sub-gradients $\partial f(x)$ for $x \in \mathbb{R}$ of the function $f : \mathbb{R} \to \mathbb{R}_+$ with $f(x) = \max(0, 1 - x)$. Show that the vector v with

$$v = \begin{cases} 2\lambda w, & yw^\top x > 1 \\ 2\lambda w, & yw^\top x = 1 \\ -yx + 2\lambda w, & yw^\top x < 1 \end{cases}$$

is $v \in \partial_w \ell(w, z = (x, y))$, aka a sub-gradient of $\ell(w, z = (x, y))$ at w, for any $w \in \mathbb{R}^d$.

(4) Write down the algorithm of online AdaGrad (Adaptive Stochastic Gradient Descent) with learning rate $\eta_t > 0$, batch size m, and termination criterion $t > T_{\max}$ for some $T_{\max} > 0$ in order to discover w^* such as

$$w^* = \arg \min_{\forall w : h_w \in \mathcal{H}} \left(\mathrm{E}_{z \sim g}\left(\ell(w, z = (x, y))\right)\right) \tag{0.7}$$

The formulas in your algorithm shoud be implemented for the above learning problem and tailored to 0.3, 0.5, and 0.6.

(5) Use the R code given below in order to generate the dataset of observed examples $S_n = \{z_i = (x_i, y_i)\}_{i=1}^n$ that contains $n = 10^6$ examples with inputs x of dimension $d = 2$. Consider $\lambda = 0$. Use a seed $w^{(0)} = (0, 0)^\top$.

(a) By using appropriate values for m, η_t and $T_{\max}$, code in R the algorithm you designed in part 4, and run it.

(b) Plot the trace plots for each of the dimensions of the generated chain $\left\{w^{(t)}\right\}$ against the iteration t.

(c) Report the value of the output w^*_{adaGrad} (any type) of the algorithm as the solution to (0.7).

(d) To which cluster y (i.e., -1 or 1) $x_{\text{new}} = (1, 0)^\top$ belongs?

```
# R code. Run it before you run anything else
#
data_generating_model <- function(n,w) {
z <- rep( NaN, times=n*3 )
z <- matrix(z, nrow = n, ncol = 3)
z[,1] <- rep(1,times=n)
z[,2] <- runif(n, min = -10, max = 10)
p <- w[1]*z[,1] + w[2]*z[,2] p <- exp(p) / (1+exp(p))
z[,3] <- rbinom(n, size = 1, prob = p)
ind <- (z[,3]==0)
z[ind,3] <- -1
x <- z[,1:2]
y <- z[,3]
return(list(z=z, x=x, y=y))
}
n_obs <- 1000000
w_true <- c(-3,4)
set.seed(2023)
out <- data_generating_model(n = n_obs, w = w_true)
set.seed(0)
z_obs <- out$z #z=(x,y)
x <- out$x
y <- out$y
#z_obs2=z_obs
#z_obs2[z_obs[,3]==-1,3]=0
#w_true <- as.numeric(glm(z_obs2[,3]~ 1+ z_obs2[,2],family = "binomial"
)$coefficients)
```

Solution.

Part 3. **Support Vector Machines**

Part 4. **The kernel trick**

Part 5. **Multi-class classification**

Part 6. **Artificial Neural Networks**

Part 7. **Gaussian process regression**

Part 8. **Revision**